



# NVIDIA RAPIDS Accelerator for Apache Spark ML

GPU Accelerated Distributed ML in Spark Clusters


Speakers: Jinfeng Li and Erik Ordentlich (joint work with Bobby Wang and Lee Yang)

The background of the slide is a dark, almost black, field. It is populated with a network of thin, light green lines that crisscross the frame. At various points where these lines intersect or terminate, there are small, bright green circular dots. Some of these dots have a slight glow or halo effect. The overall impression is one of a complex, interconnected system or network, possibly representing a neural network or a data structure.

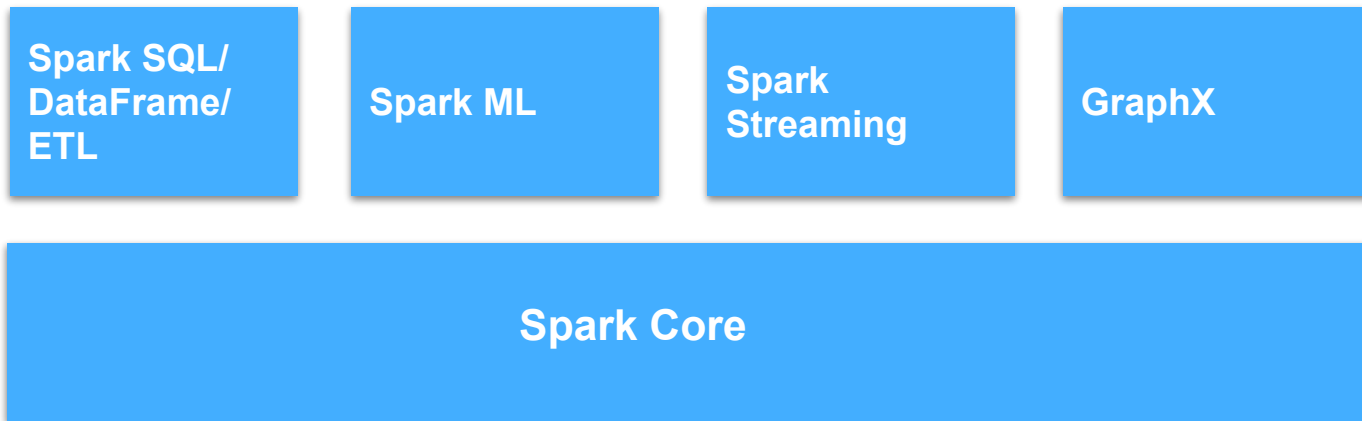
# Introduction and Motivation

# Scaling Apache Spark with GPUs

Growth in requirement  
for data processing



# Apache Spark Components



# Apache Spark Components

## GPU Acceleration


Spark SQL/  
DataFrame/  
ETL

Spark ML

# RAPIDS Accelerator for Apache Spark

## Spark ETL


- <https://github.com/NVIDIA/spark-rapids>
- Provides GPU acceleration for Spark SQL + DataFrames
- NO code change



5.7x speedup



4.5x cost savings



5x more efficient

Additional info:

<https://venturebeat.com/data-infrastructure/gtc-2023-nvidia-shares-how-rapids-can-future-proof-apache-spark>

<https://www.nvidia.com/en-us/on-demand/session/gtc-spring23-s52202/>

# RAPIDS Accelerator for Apache Spark

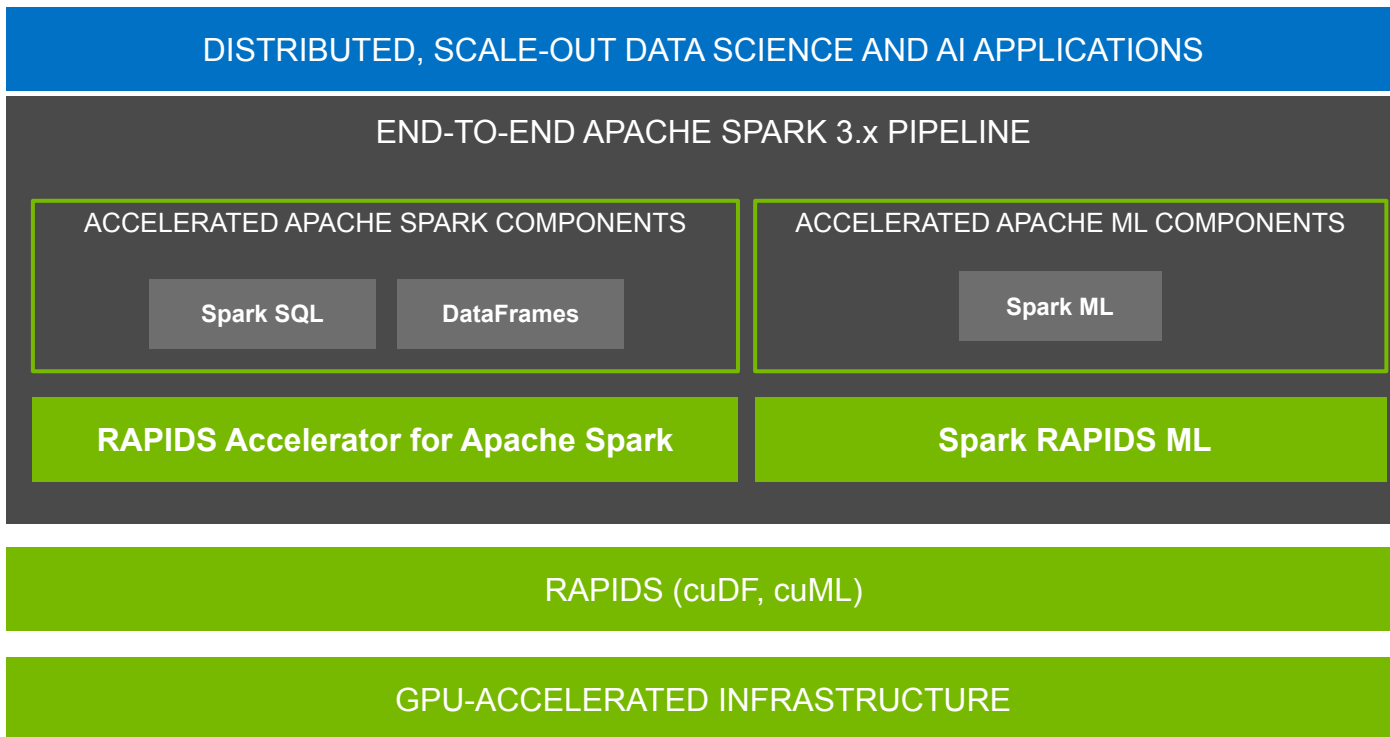
## Spark ML

- Spark ML is a key part of Apache Spark, providing distributed implementations of many ML algorithms, but **CPU only**.

### Spark RAPIDS ML

- <https://github.com/NVIDIA/spark-rapids-ml>
- New pure python open source library to GPU accelerate pySpark ML DataFrame API

# RAPIDS Accelerator for Spark ETL & ML





# Spark RAPIDS ML

## Initially Supported Algorithms

- KMeans
- PCA
- LinearRegression
  - Ridge
  - ElasticNet
- RandomForestClassifier
- RandomForestRegressor
  
- Non-Spark ML algorithm:
  - Exact K-Nearest Neighbors
    - similar APIs as Spark ML's LSH (BucketedRandomProjectionLSH)

# Spark RAPIDS ML

## Key Objectives

- **API:**
  - Compatible with [pyspark.ml DataFrame style apis](#)
  - Requires no application code change
    - Just a package import change
- **Speedup and cost benefits:**
  - Significantly improve on PySpark CPU perf and cost
- **Architecture:**
  - Leverages NVIDIA RAPIDS cuML accelerated ML library



**API**

# pyspark.ml

## pyspark.ml

```
from pyspark.ml.clustering import KMeans

kmeans_estm = KMeans()\
    .setK(100)\
    .setFeaturesCol("features")\
    .setMaxIter(30)

kmeans_model = kmeans_estm.fit(pyspark_data_frame)

kmeans_model.write().save("saved-model")

transformed = kmeans_model.transform(pyspark_data_frame)
```

# spark\_rapids\_ml

## spark\_rapids\_ml

```
from spark_rapids_ml.clustering import KMeans

kmeans_estm = KMeans()\
    .setK(100)\
    .setFeaturesCol("features")\
    .setMaxIter(30)

kmeans_model = kmeans_estm.fit(pyspark_data_frame)

kmeans_model.write().save("saved-model")

transformed = kmeans_model.transform(pyspark_data_frame)
```

# Benefits

- Spark devs do not need to learn new API
- Easy to migrate existing applications
  - Example: GPU acceleration compatible with Spark ML Pipelines, Crossvalidator, etc.

The background is a solid black field. Overlaid on this are numerous thin, light green lines that crisscross the frame in various directions, creating a complex web-like pattern. Interspersed among these lines are several small, bright green circular dots of varying sizes. Some dots appear as sharp points of light, while others have a soft, out-of-focus glow. The overall effect is reminiscent of a network diagram or a stylized representation of data flow or connectivity.

**Performance**

# Benchmarking

## Workload & Data

- estimator.fit(data\_df) [i.e. training]
  - data\_df read from Parquet format in AWS S3
- Compute intensive synthetic workloads:
  - 1 million rows
  - 3000 dimensional vectors
  - Data available in S3 public bucket.

## Environment

- Databricks AWS hosted Spark
- 3 node clusters (2 exec, 1 driver)
  - CPU clusters - 32 GB, 8 cores [m4dn.2xlarge]
  - GPU clusters - 32 GB, 8 cores, 1 NVIDIA A10 GPU [g5dn.2xlarge]

Instructions and scripts to reproduce:


<https://github.com/NVIDIA/spark-rapids-ml/tree/main/python/benchmark#databricks>

[Repo also has instructions for GCP Dataproc and AWS EMR]

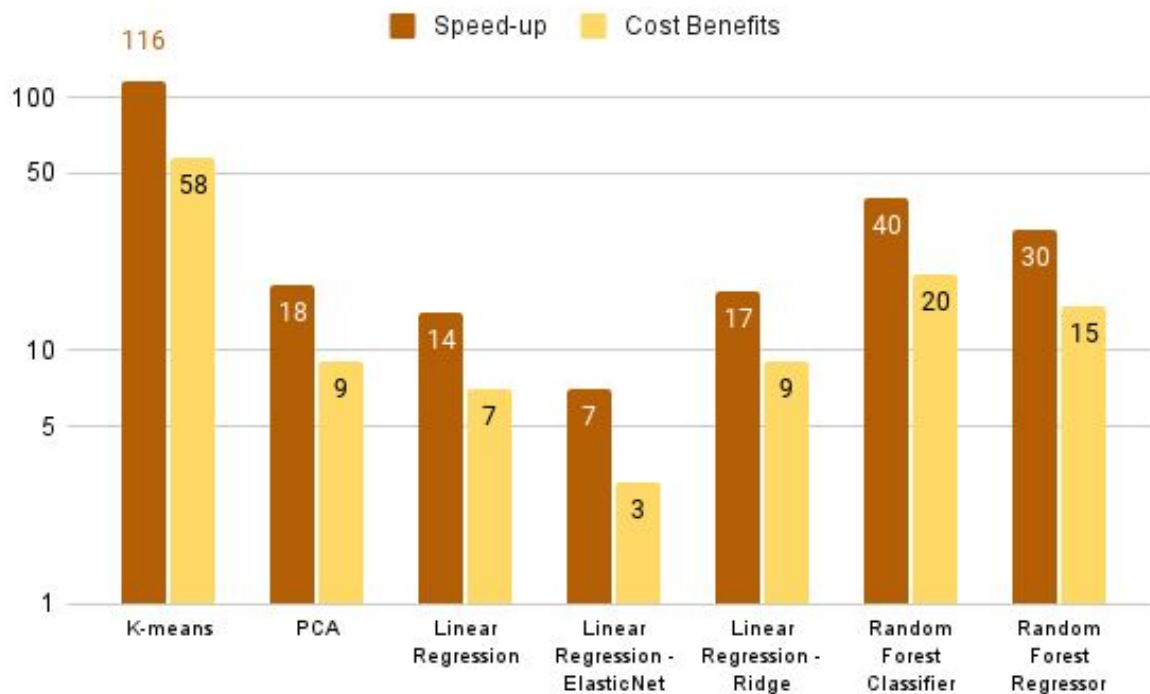


# Training/fit time: 10x-100x faster

Benchmark running times



# Cost Benefits: 3x-58x



The background is a solid black field. It is populated with a network of thin, light green lines that crisscross the frame in various directions. At several points where these lines intersect or terminate, there are small, bright green circular dots. Some of these dots have a slight glow or halo effect. The overall composition suggests a digital or technological theme, such as a network diagram or a data visualization.

**Architecture**

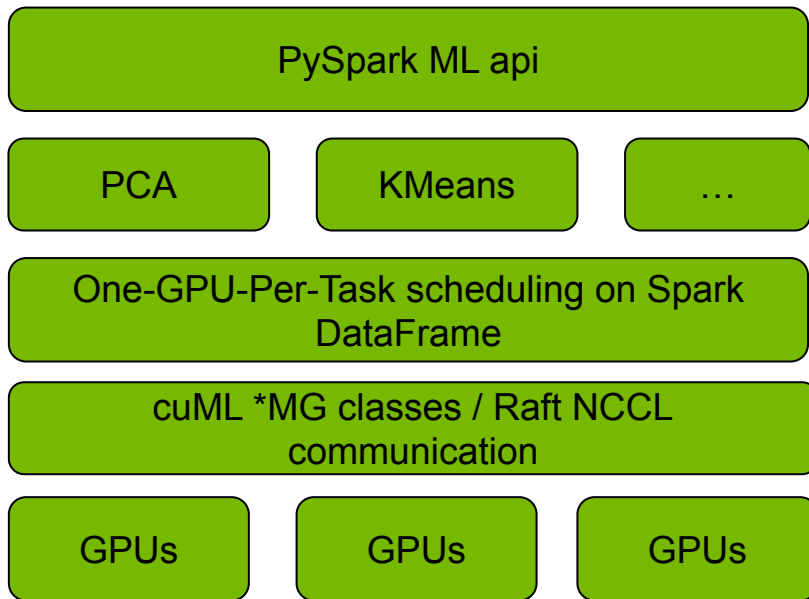
# RAPIDS cuML

- cuML is a suite of fast, GPU-accelerated machine learning algorithms
- Python API mirrors Scikit-learn, but 10-50x faster
- C++ and CUDA backend
- Diverse multi-node multi-gpu algorithms
  - GPU-accelerated communication (PCIe, NVLINK, InfiniBand Verbs)
  - DASK API

Clustering	KMeans, DBSCAN
Dimensionality Reduction	PCA, SVD, UMAP
Linear Model	Linear Regression, Lasso, Ridge, ElasticNet, Naive Bayes
Nonlinear Model	Random Forest Classifier/Regressor, KNN Classifier/Regressor
Other	NearestNeighbors

<https://github.com/rapidsai/cuml>

# Distributed cuML integration



- Use PySpark ML api for customizing algorithmic parameters and data
- Create wrappers of cuML Multi-GPU classes and raft NCCL communication layer
- Use PySpark APIs (task-per-gpu scheduling, repartition, mapInPandas, barrier, and broadcast) to setup the \*MG cluster with NCCL and process data

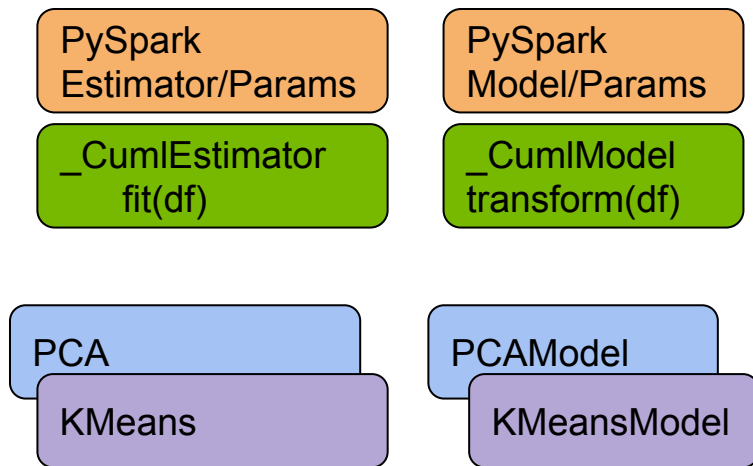
# Distributed cuML integration

```
class KMeans:
    □
    def fit(dataset):
        cuml_params = map_to_cuml(spark_ml_params)
        model = ( dataset.repartition(num_workers)
                  .mapInPandas(train_udf, ...)
                  .rdd.barrier()
                  □
                  .collect()[0] )
        return model
    □
```

# Distributed cuML integration

```
def train_udf(pandas_dfs):  
    numpy_arrays = convert(pandas_dfs)  
    comms_handle = bootstrap_comms(pyspark.BarrierTaskContext)  
    model = ( cuml....KMeansMG(cuml_params, comms_handle)  
              .fit(numpy_arrays) )  
  
    if worker==0:  
        return model  
    else:  
        return None
```

# Class architecture



- Extend the PySpark Estimator/Model to integrate into PySpark ML (eg, Pipeline, tuning)
- Automatically map PySpark algorithm parameters to cuML parameters.
- Support estimator/model persistence as in PySpark.

- The “fit” in `_CumlEstimator` handles the common steps like Data preparation, Parameter Validation, Gpu resources, cuML Context including Handle and NCCL, and Model construction and so on.
- The “transform” in `_CumlModel` similarly handles the corresponding common “transform” steps.



The background of the slide is a dark, almost black, field. It is populated with a network of thin, light green lines that crisscross the frame. At various points where these lines intersect or terminate, there are small, bright green circular dots. Some of these dots have a slight glow or halo effect. The overall impression is one of a complex, interconnected system or network, possibly representing data flow or a conceptual framework.

**Usage**

# Installation

## (Executor and Driver Environments)

- RAPIDS dependencies

```
conda create -n rapids-23.04 -c rapidsai -c nvidia -c conda-forge \
cuml=23.04 python=3.8 cudatoolkit=11.x
```

OR

```
pip install cudf-cu11 cuml-cu11 --extra-index-url=https://pypi.nvidia.com
```

- Spark RAPIDS ML

```
[ conda activate rapids-23.04 ]
pip install spark-rapids-ml
```

# Submitting Application

```
spark-submit --master ${SPARK_MASTER} --deploy-mode cluster --num-executors 2
--executor-memory 20g --driver-memory 10g \
    --conf spark.task.resource.gpu.amount=1 \
    --conf spark.executor.resource.gpu.amount=1 \
    --conf
spark.executor.resource.gpu.discoveryScript=./getGpusResources.sh \
    --conf
spark.files=${SPARK_HOME}/examples/src/main/scripts/getGpusResources.sh \
    spark-rapids-ml-kmeans-application.py
```

<https://eordentlich.github.io/spark-rapids-ml/notebooks/kmeans-with-output-databricks.html>

**Demo**

# Take aways


- Open source with Apache v2 license
- No application code change
- 3x to 50x cost benefits
- Can be run on-prem and on all CSPs



Apache Spark 3.x



Databricks



Google Cloud  
Dataproc



Amazon EMR

# Links

1.



<https://github.com/NVIDIA/spark-rapids-ml>



Tech Blog